

1. Hadoop component
2. Hadoop common
3. HDFS
4. MapReduce
5. Hadoop tools
6. Other open source tools.

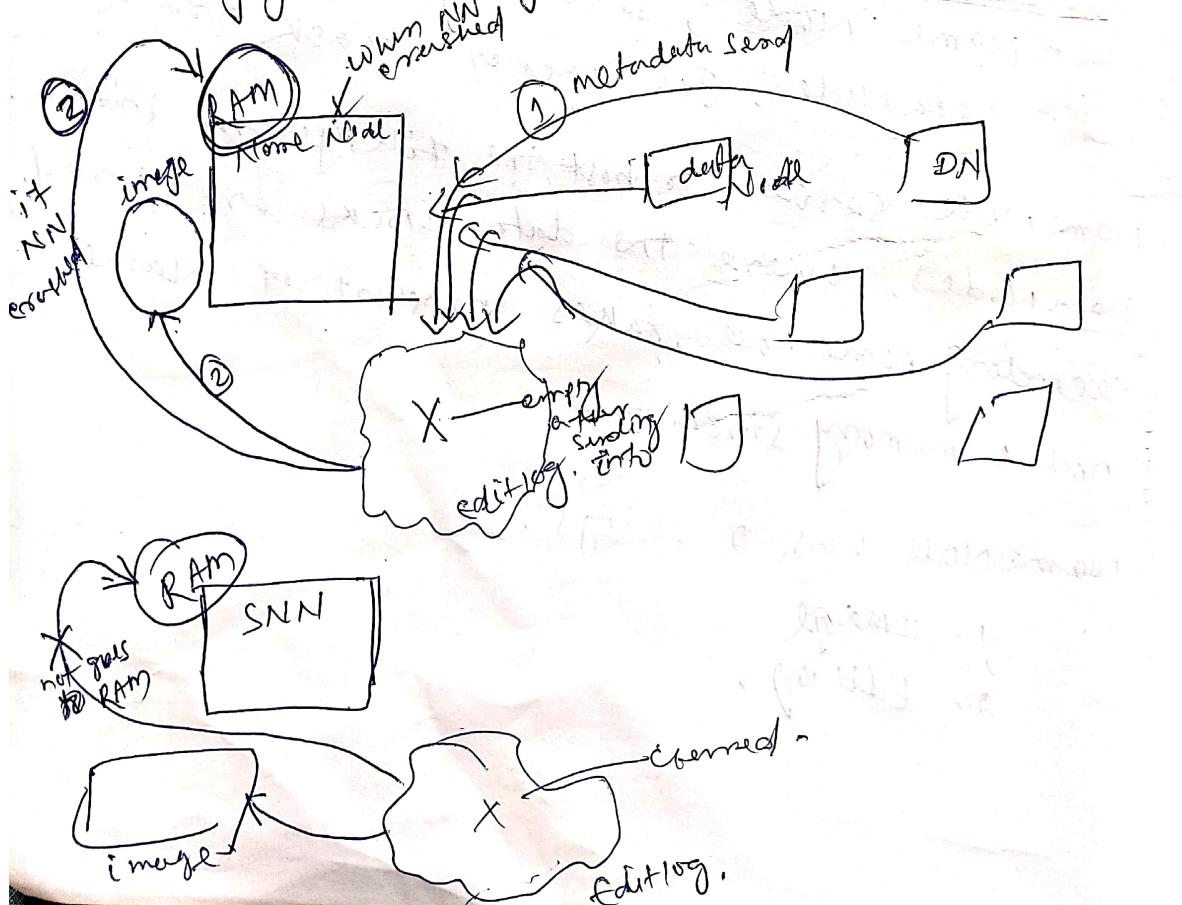
HDFS (Hadoop distributed file system)

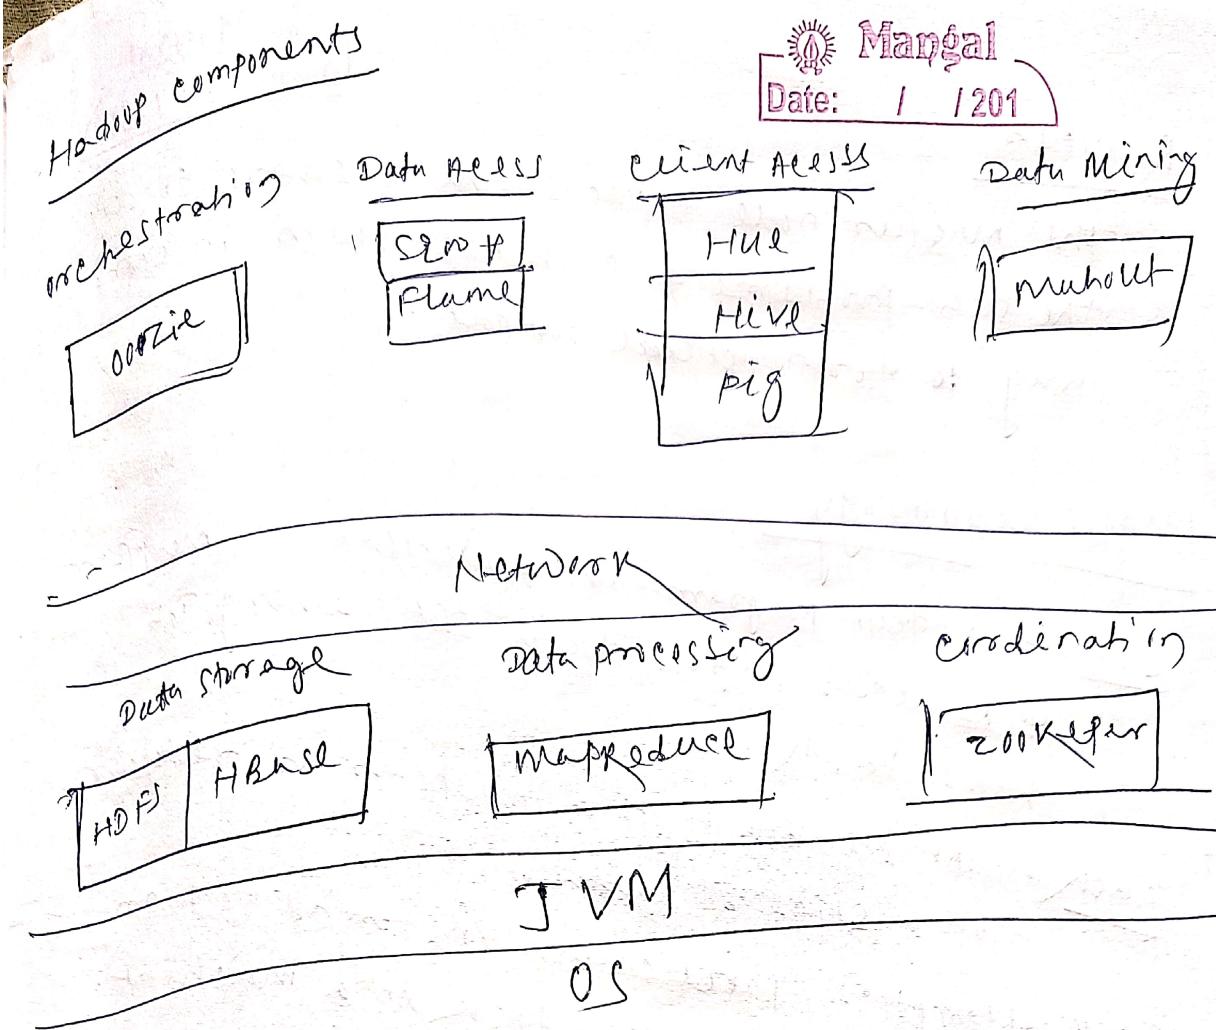
- primary storage system for Hadoop.
- distributed, portable, scalable file system
- written in Java.
- files are broken down & stored in multiple machine.
- designed for large scale distributed data processing.
- It follows Master & Slave architecture.

HDFS consists of

- Name Node, Secondary Name Node
- DataNode, Job Tracker, Task Tracker.
- NameNode server to host the filesystem index.
- DataNodes, where the data blocks are stored.
- Secondary name node, takes snapshot of Name nodes memory structure.
- NameNode has 2 things.
 - 1 - Image
 - 2 - Edit Log

- Whenever we write file to the machine, the slave node sending their info.
- Editing comes to picard. Editing have all the info?
- Let's say Namenode crashed, when coming back the editlog send the metadata to the RAM.
- When we are resending, the editlog goes to RAM as well as image.
- Then the editlog has cleared.
- When all the slave node sending metadata to the Namenode, it is commit log of secondary namenode.
- But for SNN editlog can't go to RAM. it only goes to image.



MapReduce

- A SW framework design for processing large volume of data.
- divides data into set of implementation tasks to run in parallel machine.
- It will operates in 2 steps.
 - Map step
 - Reduce step.

Mapper

- Master node takes ^{input then} partitions it into sub-problems & distributes it to worker nodes.
- Map function takes a series of key / value pairs, processes & generates 0 or more output key / value pairs.

Reducers

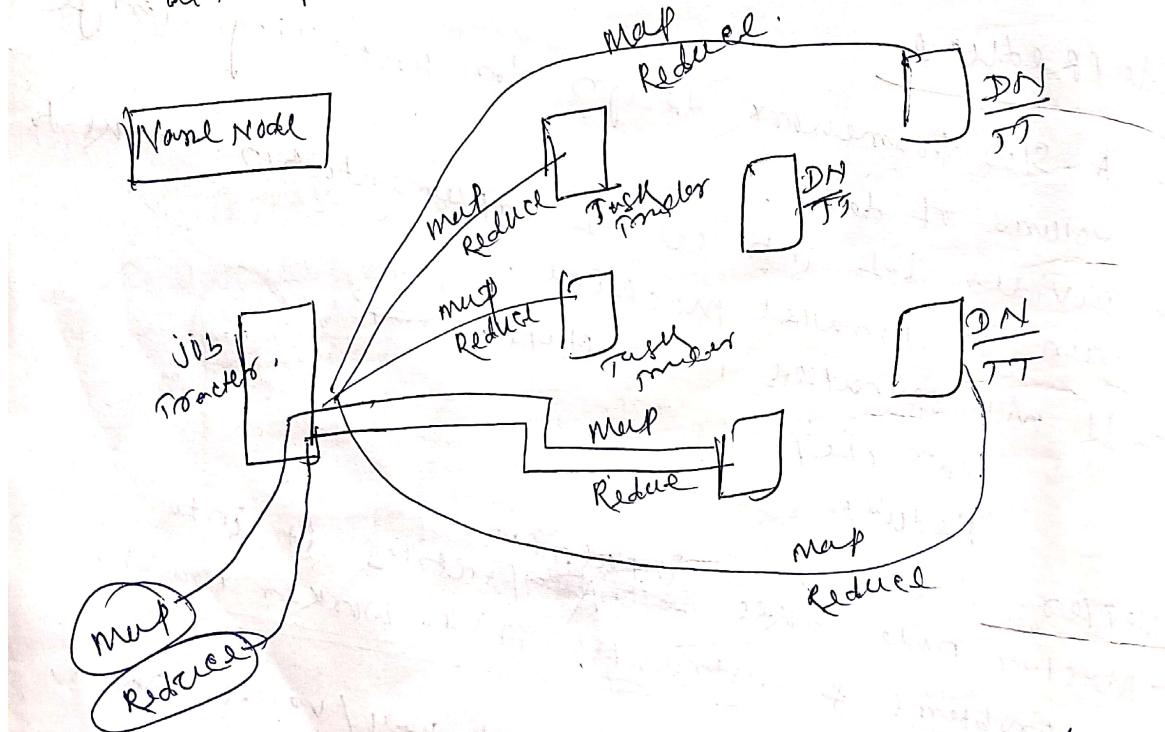
- the master node collects the answers to all the sub-problems & combines them in some way to form actual output.

Hadoop Terminology

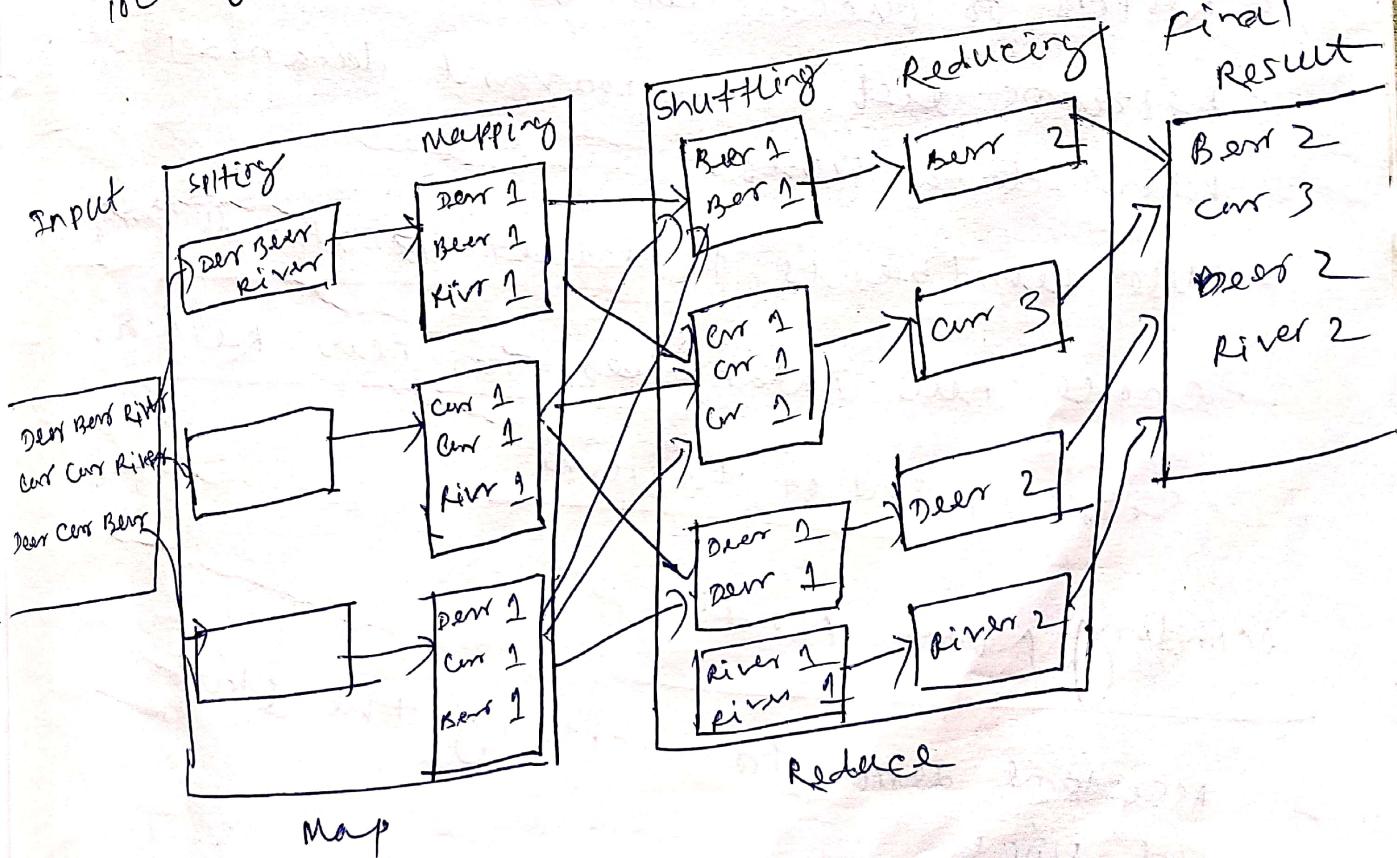
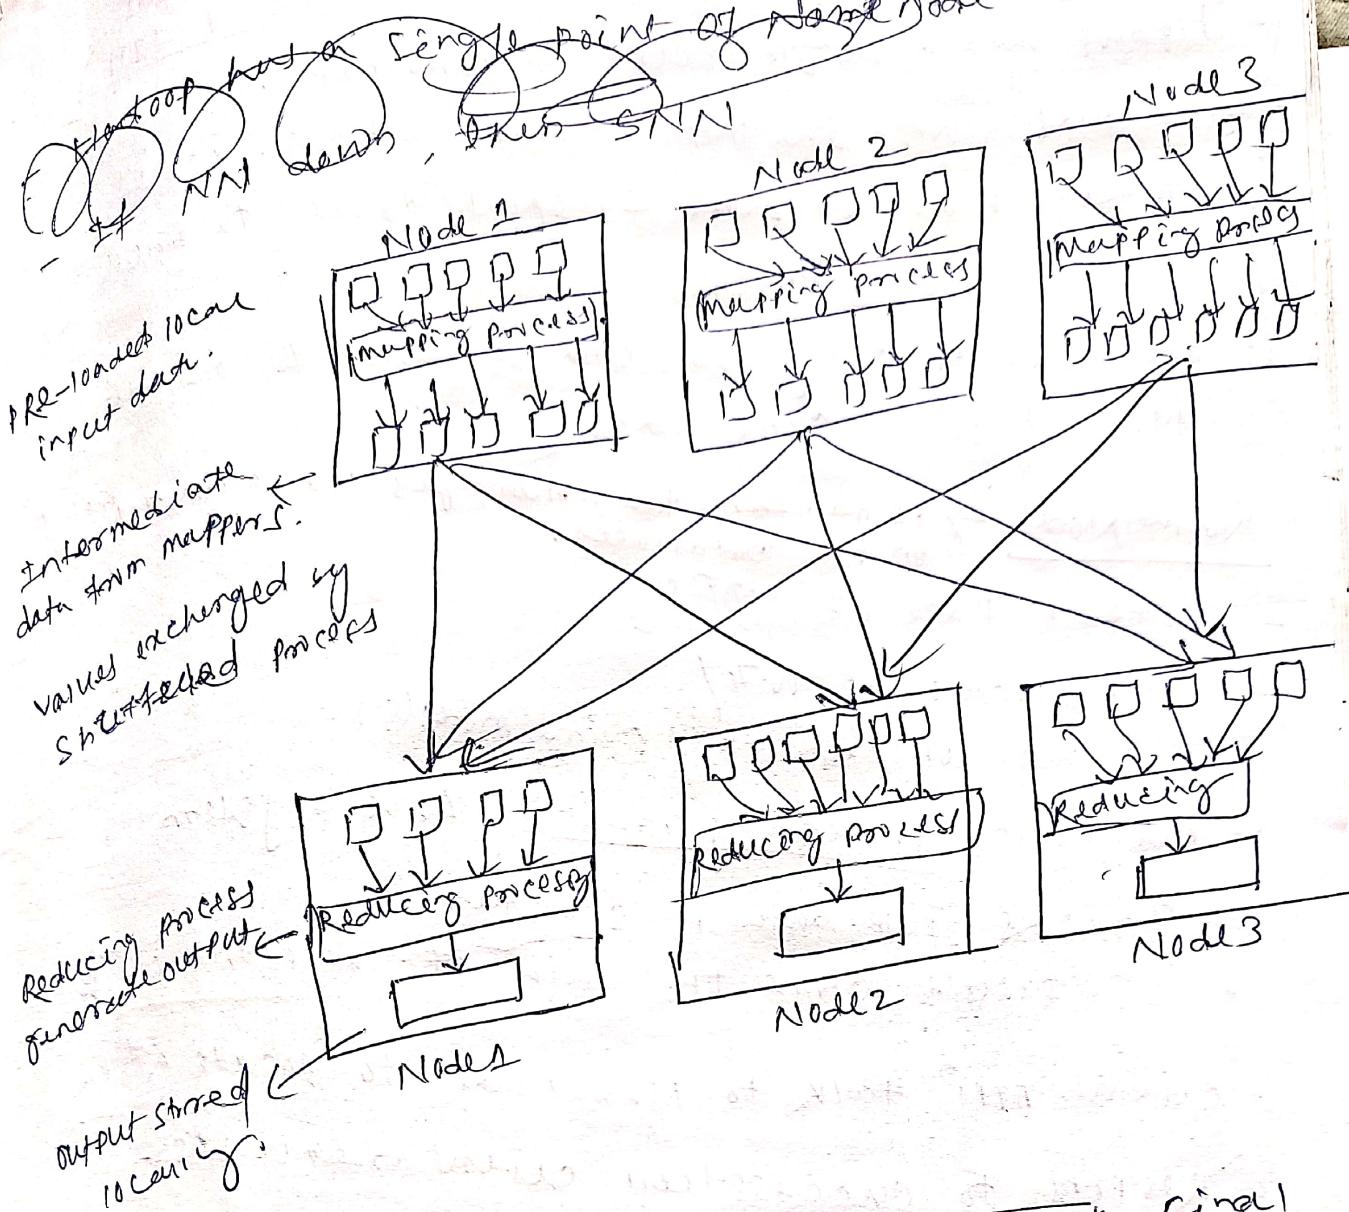
Job :- A full program - an execution of mapper & reducer across data set. every Job have task.

Task :- An execution of a mapper or reducer on a slice of data.

Task attempt :- particular instance of an attempt to execute a task on a machine.



- Job Tracker will take the input of a file & talk to task Tracker.
- Map & Reduce go to the every machine with the help of Job Tracker.



- Hadoop has a single point of name node.
 - If NN down, the SWN come back. It is taking more of time if SNN comes to back. It is a manual process.
 - We have to copy the image file for edition, then it will come back.
- NameNode → It maintains ~~table~~ & manages the ~~state representation~~ on the DataNodes.
- Center piece of HDFS
 - `http://localhost:50070/`
 - Web UI of the HDFS name node(s)
 - Maintains directory tree in the system
 - Tracks file data in cluster.
 - Doesn't store the data of these files itself.
 - Client API talk to NameNode to locate a file.
 - Response to successive client requests.
 - Returns list of relevant datanodes where data lives.
 - Use heartbeats to detect ~~dead~~ Datanode failure.
 - Choose new Datanodes for new replica.
 - Single point of failure.

Secondary Name Node

- Assistant daemon for monitoring the state of HDFS

- Each cluster has one NN residing on its own machine.
- No other datanode or TaskTracker runs in SNN.
- Doesn't receive any real-time changes to HDFS.
- It is a backup for Namenode.
- Take snapshot of metadata at intervals defined by the cluster configuration.
- Snapshots minimize the downtime & loss of data.

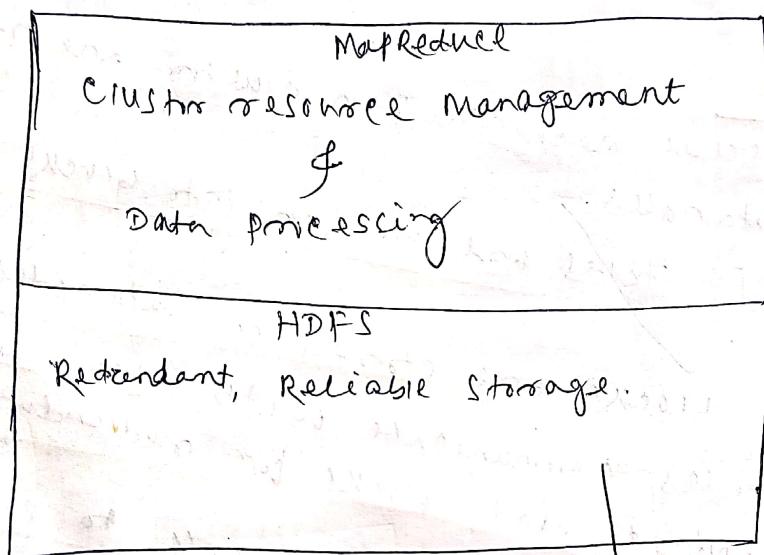
DataNode

Individual machines in a cluster are referred to as DataNodes.

- In HDFS files are broken into blocks of fixed size.
- These blocks are stored in DataNodes.
- DataNodes communicate with other DataNodes to replicate data blocks for redundancy.
- DataNodes connect to Namenode to establish service.
- Responds to requests from the Namenode.
- DataNodes send heartbeat to the Namenode.

Problem in Hadoop (1.0)

- Hadoop has a single point of NameNode.
- If NameNode down, then SNN come back. It is taking hours of time & SNN come back is a manual process.
- We have to copy the image file & editing then it will come back.



↓ which is responsible for storing the data.

- But on top of the HDFS only 1 technique is there, i.e., MapReduce.
- MapReduce is done using the help of JobTracker.
- When we submit a job, then Job Tracker will do the cluster Resource management.
↓
which machine have how many CPU & how many CPUs give to which Job.

- MapReduce doesn't have capacity
of 1000 jobs, more CPU, more RAM, more will
cause to run Job Tracker can't take more
than that.
That's why Hadoop(2.0) came to picture.

~~MapReduce~~

- single job tracker receive
all the jobs as parallel.

Apache Hadoop YARN (Hadoop 2)

- problem with large scaling.
- > 4000 nodes
- > 40K concurrent tasks.
- problem with resource utilization.
- suits only for Map or Reduce.
- single NameNode, single point of failure.
- clients & cluster must be at same version.

Apache Hadoop YARN (Hadoop 2)

- cluster resource management by Job Tracker
in MapReduce i.e., taken care by YARN.

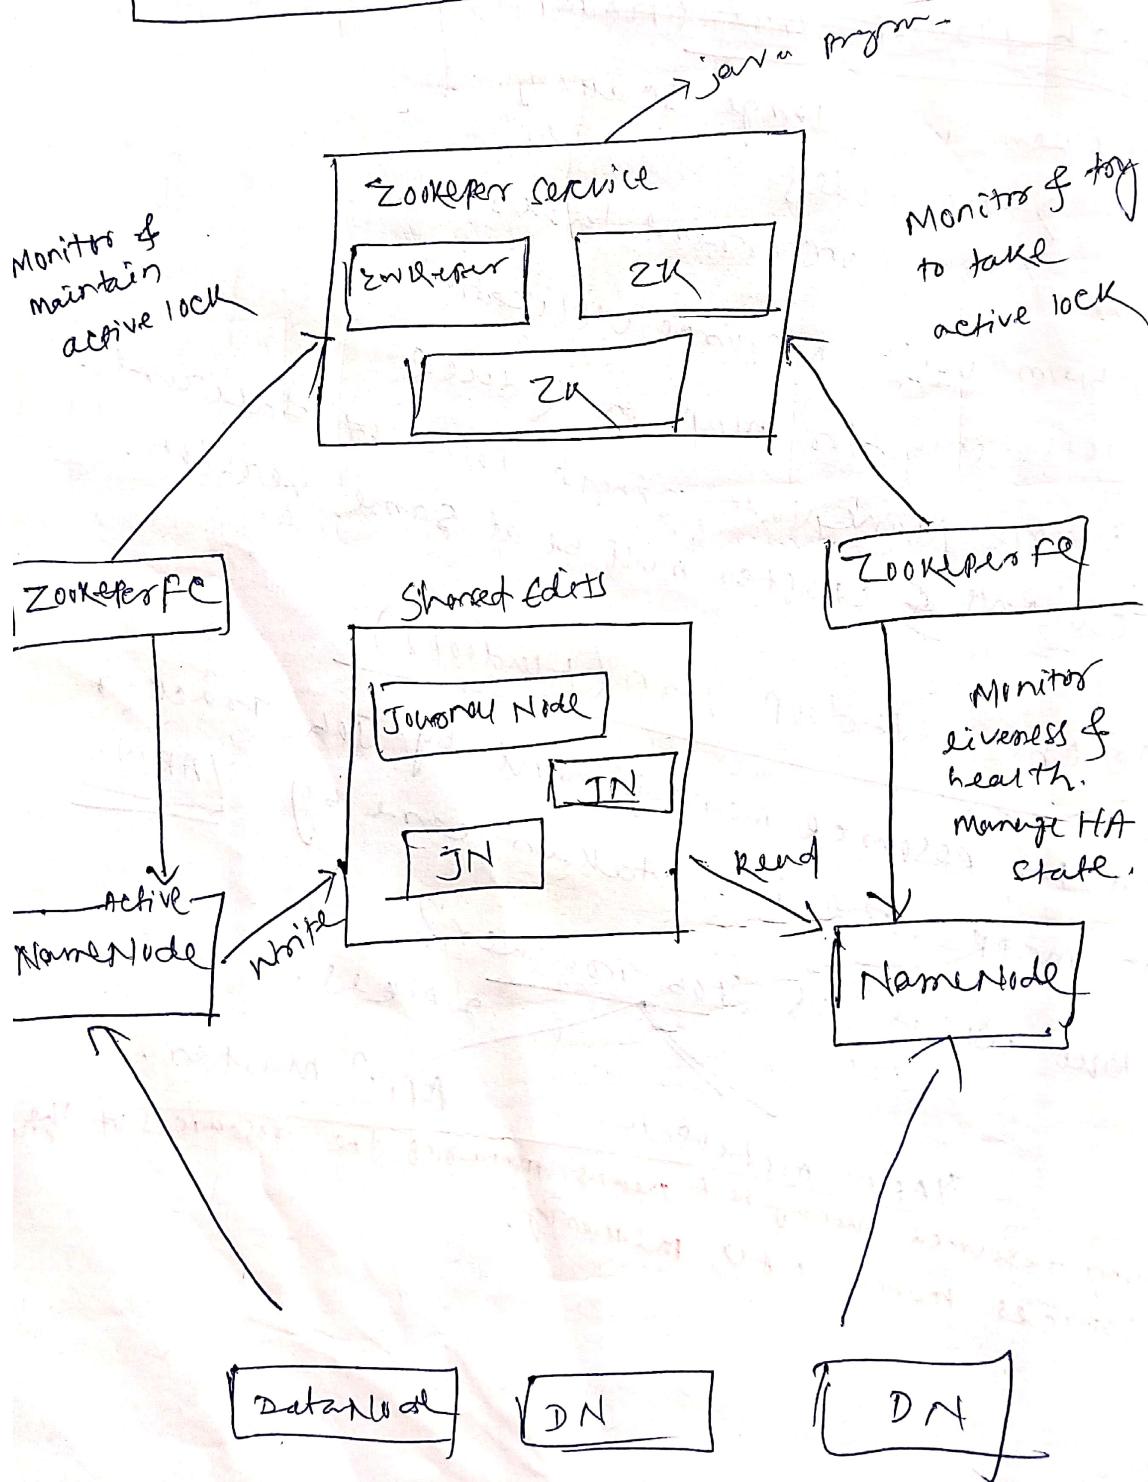
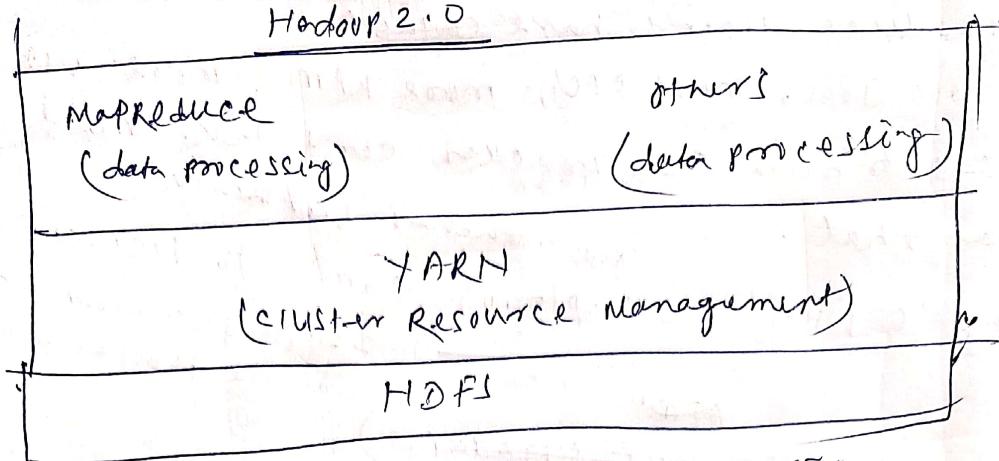
- But where is Job Tracker?

2 pieces.

YARN architecture

AppMaster

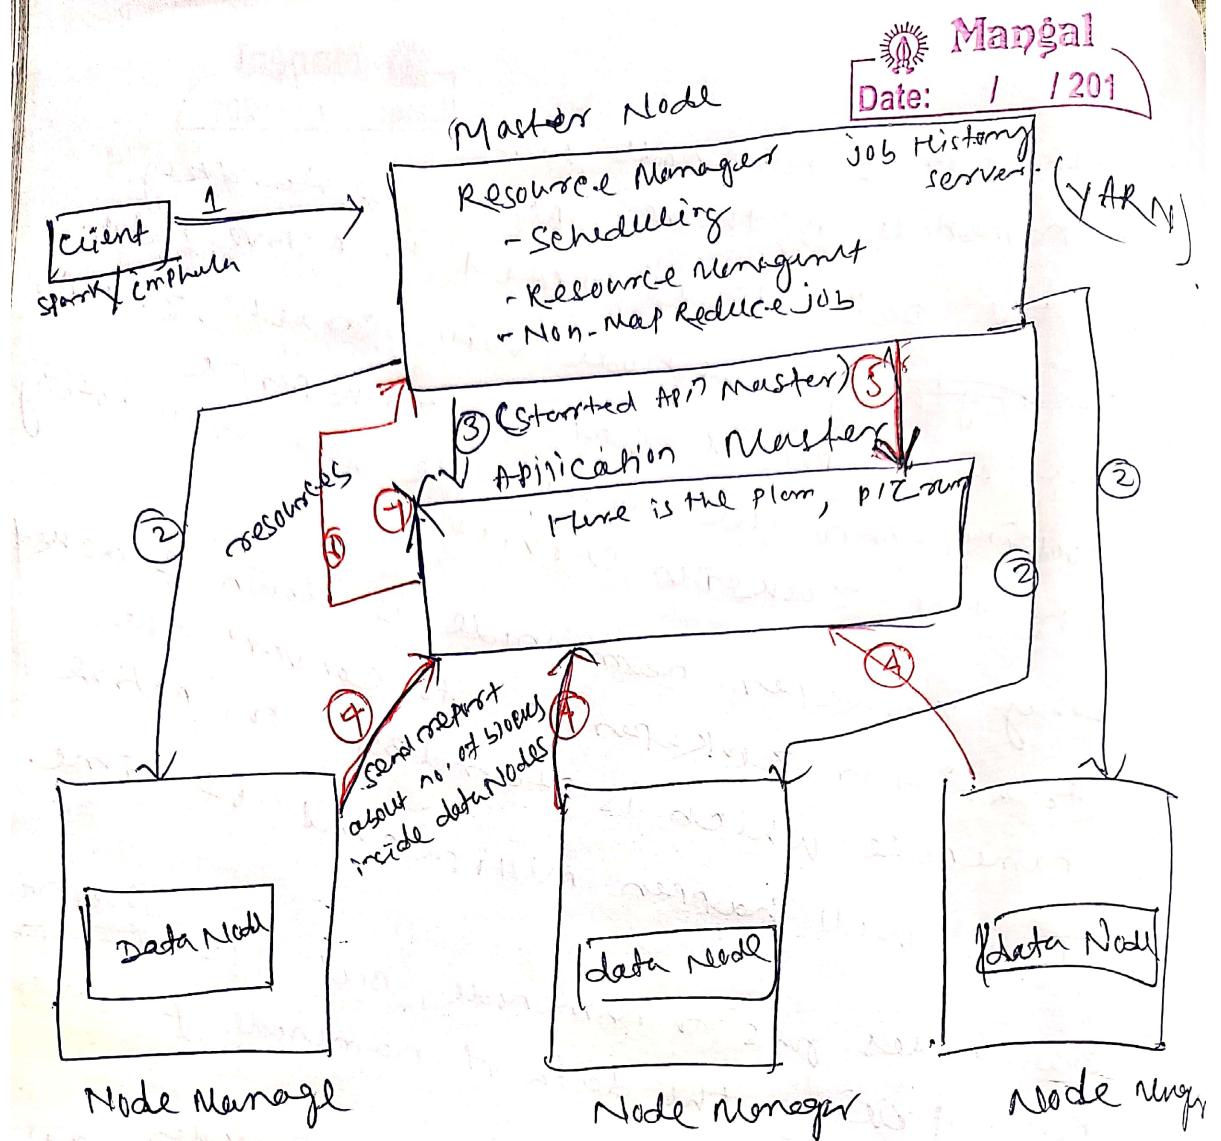
cluster resource management means managing the resources of the cluster.
resources means CPU, memory.



- Instead of 2 namenode here 1 namenode is there. Active & stand by.
- All the datanode connected to Active node.
- If ~~anything~~ name node will crash, Zookeeper
- If ~~anything~~ component called zookeeper has a component called zkkeeper looking into the name node, it will notify to the Zookeeper service.
- Hey Zookeeper name node is down & asked to another zkkeeper to convert the namenode which is ~~standby~~ to active.
- This will happen millisecond of time.
- Data nodes goes to namenode, but the ~~all~~ nodes feeding the data of namenode & converting it to ~~standby~~.

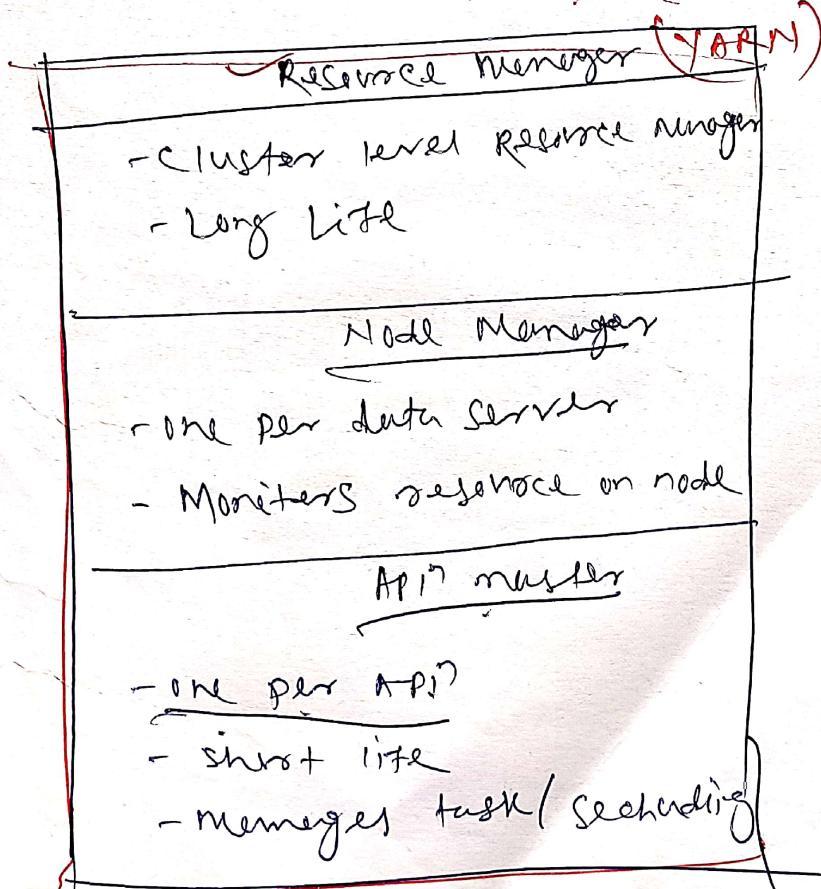
~~task tracker~~ all

- Journal node, Zookeeper, ~~datanode~~ are service only.
- Zookeeper is a java program, it runs on machine.
- Journal node taking the commit logs from Active Namenode & writing to ~~standby~~ Namenode & make as ~~shink~~.
- Task tracker replace by Node manager.
- improves scaling.
- II resource management



- client can send different type of apollo? Impala, Sparkle, ...
- non-mapreduce job will handle by Resource Manager.
- Resource manager take the data from different type of API.
- Resource Manager know how many resources are there.
- Resource Manager talk to Node Manager & then started the app master.

- API master will do app monitoring & job scheduling.
- Node manager will give the report about the no. of blocks available in the data node.
- The Resource Manager tells to API master my, API master here is the plan, PIZ is run or. (process function)
- After processing the API master returns the resources to Resource Manager.
- Then API master goes to the zone.
- There are more than 1 resource manager.



- When we submit a job the YARN will start API master.
- Then API master will start the certain no. of executors.